1 Fitting a linear model in R

Now, we shall see how to fit a linear model in our spring example. It is simple to do so in R, as follows :

```
fit=lm(length ~weight,dat)
fit
model.matrix(fit)
plot(length~weight,dat)
names(fit)
fit$coef
fit$res
fit$fit
fit$rank
abline(fit$coef)
```

Let us analyse and understand each command.

1.1 fit=lm(length ~weight, dat)

This command fits a least squares regression model where *length* is the response variable, *weight* is the input variable and *dat* is the data frame in which our information about length and weight for each trial is stored. R calculates the values of the unknown parameters using an algorithm and stores them in the object *fit*.

1.2 fit

Call:
lm(formula = length ~weight, data = dat)

Coefficients:
(Intercept) weight
3.283 2.010

This command displays the calculated values of the unknown parameters of our linear model. In our spring example we have two unknown parameters, whose values are stored in the object *fit*. Hence, $\beta_1 = 3.283, and, \beta_2 = 2.010$.

1.3 model.matrix(fit)

This command prints the design matrix which, in our example, has 6 rows (corresponding to 6 springs) and 2 columns :
(Intercept) weight
1 1.0
1 1.5
1 2.0
1 2.5
1 3.0
1 3.5

1.4 names(fit)

[1] "coefficients" "residuals" "effects" "rank" "fitted.values" "assign" "qr"
[8] "df.residual" "xlevels" "call" "terms" "model"

This command displays the various components present in the object *fit*:

coefficients : the vector of unknown parameters
residuals : the residuals, i.e, response minus fitted values.
fitted.values : the fitted mean values.
rank : the rank of the design matrix.
df.residual : the residual degrees of freedom.

1.5 components in *fit*

fit$coef
(Intercept) weight
3.283143 2.009714
fit$res
1 2 3 4 5 6
-0.002857143 0.012285714 -0.022571429 0.022571429 -0.012285714 0.002857143
fit$fit
1 2 3 4 5 6
5.292857 6.297714 7.302571 8.307429 9.312286 10.317143
fit$rank
[1] 2

1.6 plot(length ~weight,dat)

1.7 abline(fit$coef)

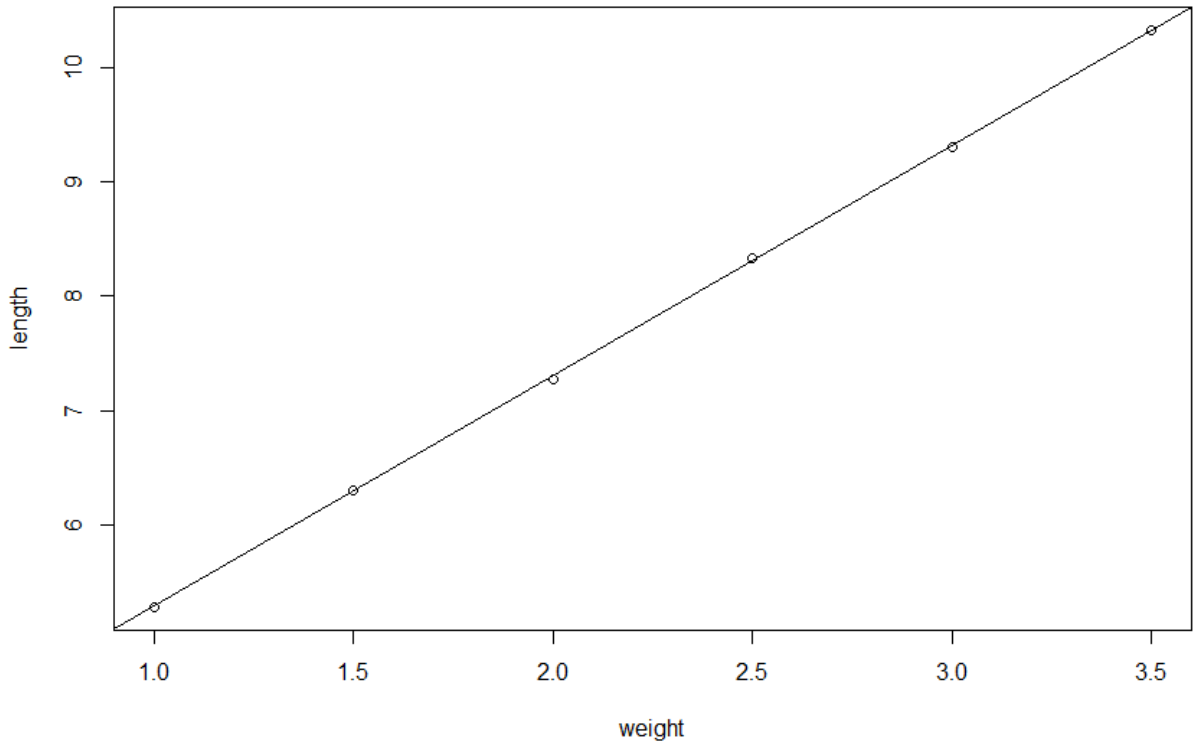

Figure 1: Fitted linear model in R.

This command displays a graph of length versus weight. The *abline()* command gives a straight line with intercept and slope as the calculated coefficients. Here we see that the 6 points are almost lying on a straight line, with hardly any random variation. Hence, a linear model would be a very good fit for our particular example.